



7/29/2022

Project-01

Configure more than one webserver with proper load balancing, auto scaling and SSL configuration.

Sumit Mishra
SIC: 190310286

1. Creating and configuring a VPC.

1. Create a VPC

	Name	VPC ID	State	IPv4 CIDR	IPv6 CIDR	DHCP
	personal-vpc	vpc-083ef3fe083e0f251	Available	7.0.0.0/16	-	dopt-0

2. Create a subnet inside the VPC to host servers.

subnet-0c89d2c7eab4ba5ac / web-subnet Actions

Details

Subnet ID subnet-0c89d2c7eab4ba5ac	Subnet ARN arn:aws:ec2:ap-south-1:258046353232:subnet/subnet-0c89d2c7eab4ba5ac	State Available	IPv4 CIDR 7.0.1.0/24
Available IPv4 addresses 247	IPv6 CIDR -	Availability Zone ap-south-1a	Availability Zone ID aps1-az1
VPC vpc-083ef3fe083e0f251 personal-vpc	Route table rtb-0b7cda26fa0da090d web-server-rt	Network ACL acl-000ec8cc218e077b7	Default subnet No
Auto-assign public IPv4 address Yes	Auto-assign IPv6 address No	Auto-assign customer-owned IPv4 address No	Customer-owned IPv4 pool -
Outpost ID -	IPv4 CIDR reservations -	IPv6 CIDR reservations -	IPv6-only No
Hostname type IP name	Resource name DNS A record Disabled	Resource name DNS AAAA record Disabled	DNS64 Disabled
Owner 258046353232			

3. Created an Internet Gateway for the servers to access the internet through it.

VPC > Internet gateways > igw-0f10d2feb4472a074

igw-0f10d2feb4472a074 / personal-igw Actions

Details Info

Internet gateway ID igw-0f10d2feb4472a074	State Attached	VPC ID vpc-083ef3fe083e0f251 personal-vpc	Owner 258046353232
--	-------------------	--	-----------------------

Tags

< **1** > ⚙

Key	Value
Name	personal-igw

4. Created and configured the route table for the subnet created.

VPC > Route tables > rtb-0b7cda26fa0da090d

rtb-0b7cda26fa0da090d / web-server-rt

Actions

You can now check network connectivity with Reachability Analyzer

Run Reachability Analyzer

DetailsInfo

Route table ID

rtb-0b7cda26fa0da090d

VPC

vpc-083ef3fe083e0f251 | personal-vpc

Main

No

Owner ID

258046353232

Explicit subnet associations

subnet-0c89d2c7eab4ba5ac / web-subnet

Edge associations

-

Routes

Subnet associations

Edge associations

Route propagation

Tags

Routes (6)

Edit routes

Filter routes

Both

< 1 >

Destination	Target	Status	Propagated
0.0.0.0/0	igw-0f10d2feb4472a074	Active	No
7.0.0.0/16	local	Active	No
13.232.122.87/32	igw-0f10d2feb4472a074	Active	No
13.233.197.113/32	igw-0f10d2feb4472a074	Active	No
43.205.130.90/32	igw-0f10d2feb4472a074	Active	No
65.0.95.54/32	igw-0f10d2feb4472a074	Active	No

2. Configuring Web servers, domain, hosting.

1. Launching 4 EC2 instances with Linux AMI.

Instances (5) Info

Refresh

Connect

Instance state

Actions

Launch instances

Search

< 1 >

	Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IPv4 DNS	Public IPv4 ...
<input type="checkbox"/>	webserver4	i-0c31d18055bbbc84d	Running	t2.micro	2/2 checks passed	No alarms	ap-south-1a	-	3.110.18.77
<input type="checkbox"/>	webserver3	i-0ca0d578d9ae4db13	Running	t2.micro	2/2 checks passed	No alarms	ap-south-1a	-	13.235.92.242
<input type="checkbox"/>	webserver1	i-01ef2264909f51ee6	Running	t2.micro	2/2 checks passed	No alarms	ap-south-1a	-	43.205.77.193
<input type="checkbox"/>	webserver2	i-0d6fe1c4744994690	Running	t2.micro	2/2 checks passed	No alarms	ap-south-1a	-	43.205.53.124

2. Assigning Elastic IP addresses to all the four servers.

Instances (5) Info

Launch instances

Search

< 1 >

	Name	Instance ID	Instance state	Instance type	Status check	Public IPv4 ...	Elastic IP
<input type="checkbox"/>	webserver4	i-0c31d18055bbbc84d	Running	t2.micro	2/2 checks passed	3.110.18.77	3.110.18.77
<input type="checkbox"/>	webserver3	i-0ca0d578d9ae4db13	Running	t2.micro	2/2 checks passed	13.235.92.242	13.235.92.242
<input type="checkbox"/>	webserver1	i-01ef2264909f51ee6	Running	t2.micro	2/2 checks passed	43.205.77.193	43.205.77.193
<input type="checkbox"/>	webserver2	i-0d6fe1c4744994690	Running	t2.micro	2/2 checks passed	43.205.53.124	43.205.53.124

3. Creating a Hosted zone in AWS Route 53 and adding hosting records for root domain(sumitmishra.info) and sub-domain(www.sumitmishra.info).

The screenshot shows the AWS Route 53 console for the hosted zone 'sumitmishra.info'. The 'Records (4)' tab is selected, displaying a table of DNS records. The records include an A record for the root domain, an NS record for nameservers, an SOA record, and an A record for the www sub-domain.

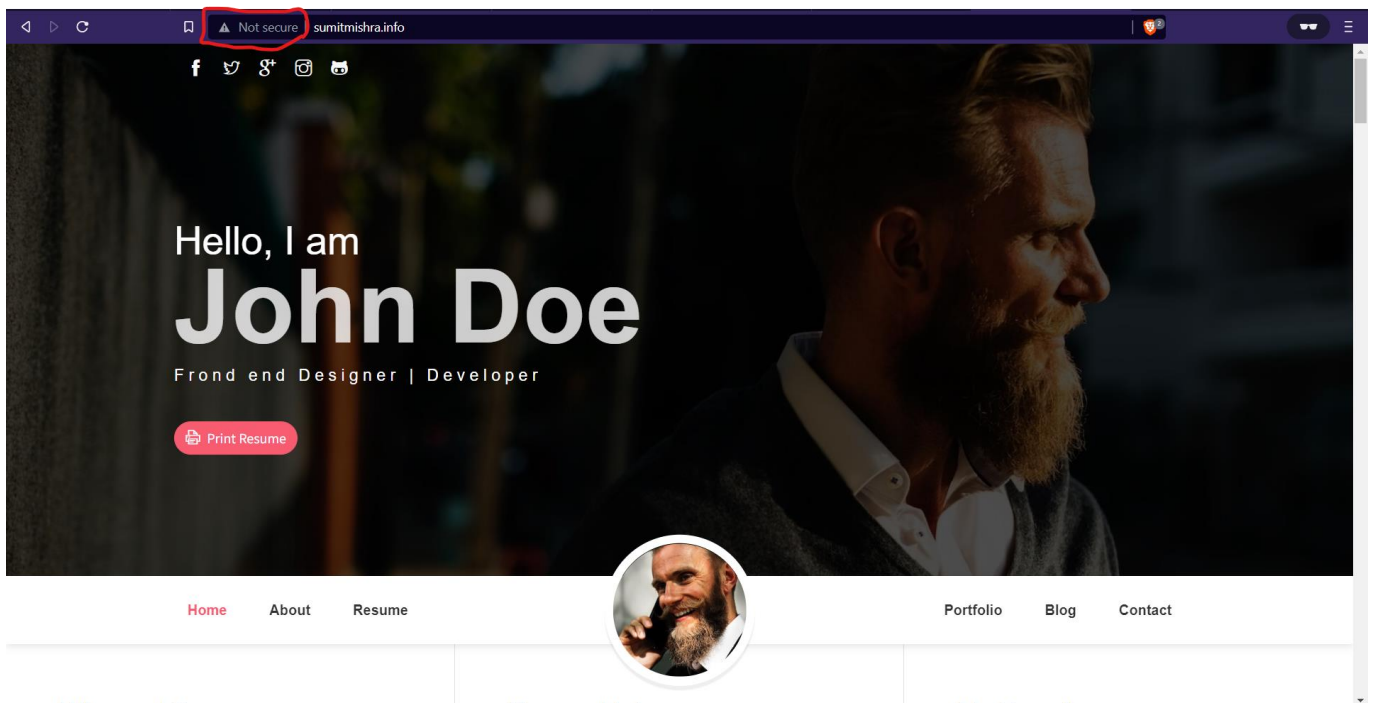
Record name	Type	Routin...	Differ...	Value/Route traffic to
sumitmishra.info	A	Simple	-	43.205.77.193
sumitmishra.info	NS	Simple	-	ns-15.awsdns-01.com. ns-1594.awsdns-07.co.uk. ns-985.awsdns-59.net. ns-1149.awsdns-15.org.
sumitmishra.info	SOA	Simple	-	ns-15.awsdns-01.com. awsdns-hostmaster.amazon.com. 1 7200 900 1209600 86400
www.sumitmishra.info	A	Simple	-	43.205.77.193

4. Replacing Nameservers in GoDaddy Domain (sumitmishra.info) with AWS Route 53 Nameservers.

The screenshot shows the GoDaddy Nameservers management page. It indicates that custom nameservers are being used and provides a 'Change' button. Below, the current nameservers are listed: ns-15.awsdns-01.com, ns-1594.awsdns-07.co.uk, ns-985.awsdns-59.net, and ns-1149.awsdns-15.org.

3. Configuring SSL/TLS certificate for a webserver using let us encrypt Apache certbot.

1. Opening the website hosted in webserver one without SSL/TLS certificate we get the not secured warning in the left most side of URL bar.



2. Now configuring the SSL certificate through let's encrypt using ssh client.

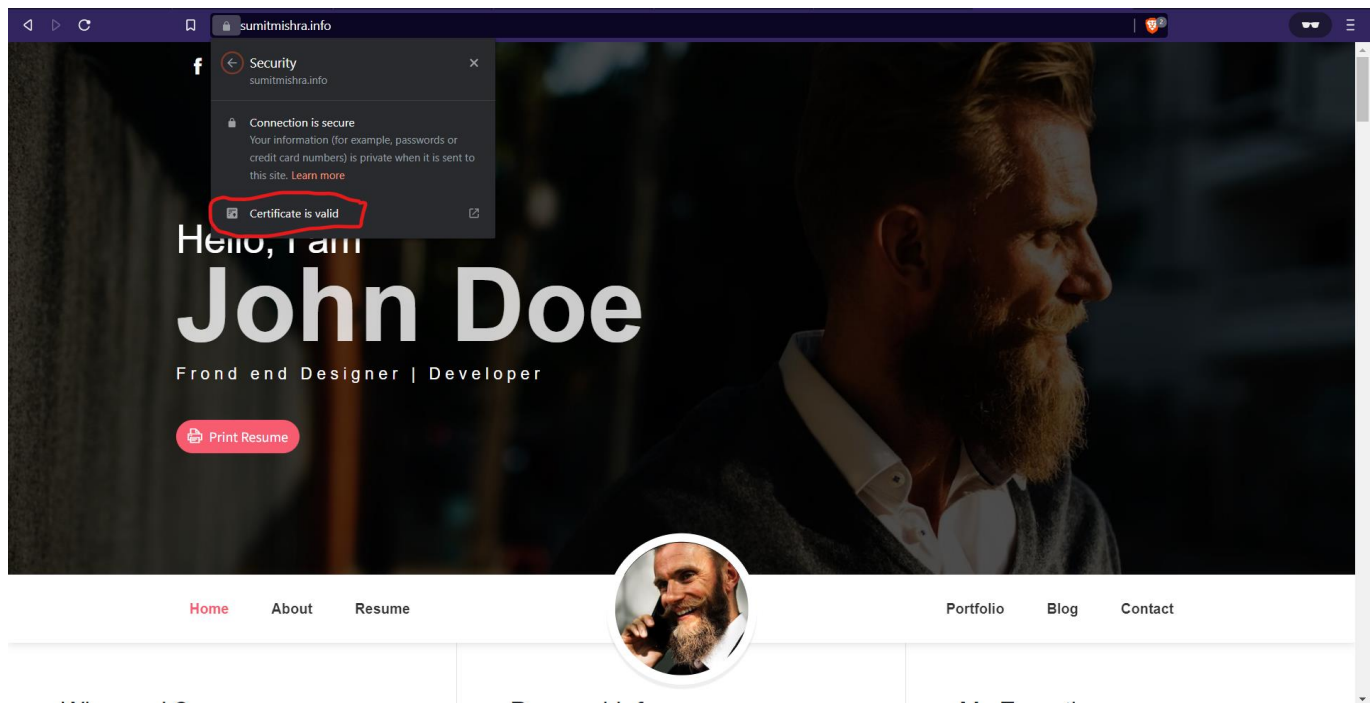
```
[root@ip-7-0-1-148 ~]# certbot --apache
Saving debug log to /var/log/letsencrypt/letsencrypt.log
Plugins selected: Authenticator apache, Installer apache

Which names would you like to activate HTTPS for?
- - - - -
1: sumitmishra.info
2: www.sumitmishra.info
- - - - -

Select the appropriate numbers separated by commas and/or spaces, or leave input
blank to select all options shown (Enter 'c' to cancel):
Requesting a certificate for sumitmishra.info and www.sumitmishra.info
Performing the following challenges:
http-01 challenge for sumitmishra.info
http-01 challenge for www.sumitmishra.info
Waiting for verification...
Cleaning up challenges
Created an SSL vhost at /etc/httpd/conf.d/vhost-le-ssl.conf
Deploying Certificate to VirtualHost /etc/httpd/conf.d/vhost-le-ssl.conf
Deploying Certificate to VirtualHost /etc/httpd/conf.d/vhost-le-ssl.conf
Redirecting vhost in /etc/httpd/conf.d/vhost.conf to ssl vhost in /etc/httpd/conf.d/vhost-le-ssl.conf

- - - - -
Congratulations! You have successfully enabled https://sumitmishra.info and
https://www.sumitmishra.info
- - - - -
```

3. Now again hitting the URL sumitmishra.info and www.sumitmishra.info to see that the website shows secured in the leftmost area of the URL bar, this means that SSL certificate has been enabled.



4. Creating a target group using the four servers and configuring load balancing along with SSL certificate using let's encrypt for the load balancer.

1. Creating a target group.

EC2 > Target groups

Target groups (1/1) Info Refresh Actions Create target group

Search or filter target groups

<input checked="" type="checkbox"/>	Name	ARN	Port	Protocol	Target type	Load balancer	VPC ID
<input checked="" type="checkbox"/>	personal-vpc-tg	arn:aws:elasticloadbalancing...	80	HTTP	Instance	None associated	vpc-083ef3fe083e0f25

Target group: personal-vpc-tg

Registered targets (4) Refresh Deregister Register targets

Filter resources by property or value

<input type="checkbox"/>	Instance ID	Name	Port	Zone	Health status	Health status details
<input type="checkbox"/>	i-0c30c9909807edd2d	webserver1	80	ap-south-1a	unused	Target group is not configured to receive traffic from the load balancer
<input type="checkbox"/>	i-02640c5161f8a7117	webserver4	80	ap-south-1b	unused	Target group is not configured to receive traffic from the load balancer
<input type="checkbox"/>	i-0bbe0940452be4101	webserver2	80	ap-south-1b	unused	Target group is not configured to receive traffic from the load balancer
<input type="checkbox"/>	i-0e0945972773e248c	webserver3	80	ap-south-1a	unused	Target group is not configured to receive traffic from the load balancer

2. Creating a load balancer.

Basic configuration

Load balancer name

Name must be unique within your AWS account and cannot be changed after the load balancer is created.

A maximum of 32 alphanumeric characters including hyphens are allowed, but the name must not begin or end with a hyphen.

Scheme [Info](#)

Scheme cannot be changed after the load balancer is created.

☒ Internet-facing

An internet-facing load balancer routes requests from clients over the internet to targets. Requires a public subnet. [Learn more](#)

☐ Internal

An internal load balancer routes requests from clients to targets using private IP addresses.

IP address type [Info](#)

Select the type of IP addresses that your subnets use.

☒ IPv4

Recommended for internal load balancers.

☐ Dualstack

Includes IPv4 and IPv6 addresses.

Network mapping [Info](#)

The load balancer routes traffic to targets in the selected subnets, and in accordance with your IP address settings.

VPC [Info](#)

Select the virtual private cloud (VPC) for your targets. Only VPCs with an internet gateway are enabled for selection. The selected VPC cannot be changed after the load balancer is created. To confirm the VPC for your targets, view your [target groups](#).

vpc-083ef3fe083e0f251

IPv4: 7.0.0.0/16



Mappings [Info](#)

Select at least one Availability Zone and one subnet for each zone. We recommend selecting at least two Availability Zones. The load balancer will route traffic only to targets in the selected Availability Zones. Zones that are not supported by the load balancer or VPC cannot be selected. Subnets can be added, but not removed, once a load balancer is created.

☒ ap-south-1a

Subnet

web-subnet-1 ▼

IPv4 settings

Assigned by AWS

☒ ap-south-1b

Subnet

web-subnet-2 ▼

IPv4 settings

Assigned by AWS

EC2 > Security Groups > sg-05a720fd534a6ba81 - personal-vpc-lb-sg

sg-05a720fd534a6ba81 - personal-vpc-lb-sg

Actions

Details

Security group name

personal-vpc-lb-sg

Security group ID

sg-05a720fd534a6ba81

Description

security group for personal VPC subnet Load balancer

VPC ID

vpc-083ef3fe083e0f251

Owner

258046353232

Inbound rules count

2 Permission entries

Outbound rules count

1 Permission entry

Inbound rules

Outbound rules

Tags

Inbound rules (2)

Filter security group rules

< 1 >

Manage tags

Edit inbound rules

	Name	Security group rule...	IP version	Type	Protocol	Port range	Source	Description
<input type="checkbox"/>	-	sgr-0515a0ed07dcd40...	IPv4	HTTP	TCP	80	0.0.0.0/0	-
<input type="checkbox"/>	-	sgr-06335d143f137fa86	IPv4	HTTPS	TCP	443	0.0.0.0/0	-

search: personal-vpc-lb

Add filter

<<

1 to 1 of 1

>>

<div><div></div></div> Name	DNS name	State	VPC ID	Availability Zones	Type	Created At	Monitoring
<div><div></div></div> personal-vpc-lb	personal-vpc-lb-845205905...	Provisioning	vpc-083ef3fe083e0f251	ap-south-1b, ap-south-1a	application	July 30, 2022 at 9:39:23 AM ...	<div><div></div></div>

Load balancer: personal-vpc-lb

Description

Listeners

Monitoring

Integrated services

Tags

Basic Configuration

Name

personal-vpc-lb

ARN

arn:aws:elasticloadbalancing:ap-south-1:258046353232:loadbalancer/app/personal-vpc-lb/03148219436b89bf

DNS name

personal-vpc-lb-845205905.ap-south-1.elb.amazonaws.com

State

Provisioning

Type

application

Scheme

Internet-facing

IP address type

ipv4

VPC

vpc-083ef3fe083e0f251

Availability Zones

subnet-0225018fd4692efec - ap-south-1b

3. Creating a new sub-domain and associating the load balancer’s DNS to that domain name.

Record details

Edit record

Record name

app.sumitmishra.info

Record type

A

Value

dualstack.personal-vpc-lb-845205905.ap-south-1.elb.amazonaws.com.

Alias

Yes

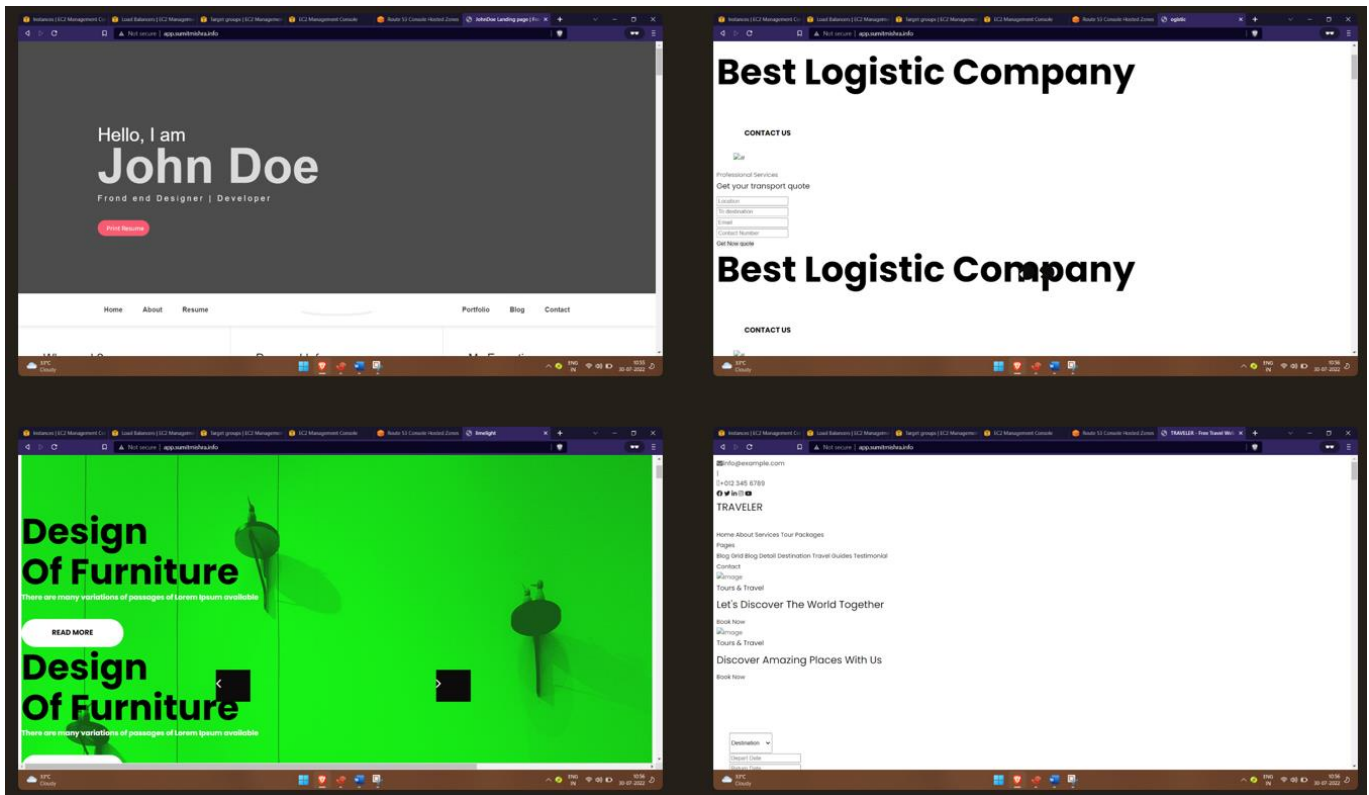
TTL (seconds)

-

Routing policy

Simple

4. Server is distributing the traffic, but SSL certificate has not been configured.



5. SSL configuration using AWS Certificate Manager.
 - a. Request a certificate.

AWS Certificate Manager > Certificates > Request certificate > Request public certificate

Request public certificate

Domain names

Fully qualified domain name [Info](#)

You can add additional names to this certificate. For example, if you're requesting a certificate for "www.example.com", you might want to add the name "example.com" so that customers can reach your site by either name.

Select validation method [Info](#)

Select a method for validating domain ownership

☒ **DNS validation - recommended**
Choose this option if you are authorized to modify the DNS configuration for the domains in your certificate request.

☐ **Email validation**
Choose this option if you do not have permission or cannot obtain permission to modify the DNS configuration for the domains in your certificate request.

Tags [Info](#)

To help you manage your certificates you can optionally assign your own metadata to each resource in the form of tags.

Tag key

Tag value - optional

Custom tag value

You can add 49 more tag(s).

- b. After requesting, add the records to the route 53 table.

Domains (2)				Create records in Route 53	Export to CSV
< 1 >					
Domain	Status	Renewal status	Type	CNAME name	CNAME value
sumitmishra.info	Success	-	CNAME	_a9ef0869ebcb98a8cef7a1082a6280e1.sumitmishra.info.	_d9b8e5f79784562a53896b12c2adeb95.vrztfgqhj.acm-validations.aws.
*.sumitmishra.info	Success	-	CNAME	_a9ef0869ebcb98a8cef7a1082a6280e1.sumitmishra.info.	_d9b8e5f79784562a53896b12c2adeb95.vrztfgqhj.acm-validations.aws.

- c. After certificate has been issued, go to load balancer -> listener -> add new listener -> select https and forward it to target group created and select the certificate as the created one.

arn:aws:elasticloadbalancing:ap-south-1:258046353232:listener/app/personal-vpc-lb/03148219436b89bf/aa0d128763...

Listener details

A listener is a process that checks for connection requests, using the protocol and port you configure. Traffic received by the listener is then routed per your specification. You can specify multiple rules and multiple certificates per listener after the load balancer is created.

Protocol: **HTTPS** : Port: **443**
1-65535

Default actions [Info](#)
Specify the default actions for traffic on this listener. Default actions apply to traffic that does not meet the conditions of rules on your listener. Rules can be configured after the listener is created.

▼ 1. Forward to [Info](#) Remove

Target group: **personal-vpc-tg** HTTP Weight (0-999): 1
Target type: Instance, IPv4
Traffic distribution: 100%

Select a target group ▼ 0 ×

[Create target group](#)

☐ [Enable group-level stickiness](#) [Info](#)
If you enable stickiness for your target group, requests routed to it remain in the same group for the duration you specify.

Add action ▼

Secure listener settings [Info](#)

Security policy
Your load balancer uses a Secure Socket Layer (SSL) negotiation configuration, known as a security policy, to negotiate SSL connections with clients.
ELBSecurityPolicy-2016-08 ▼

[Compare security policies](#)

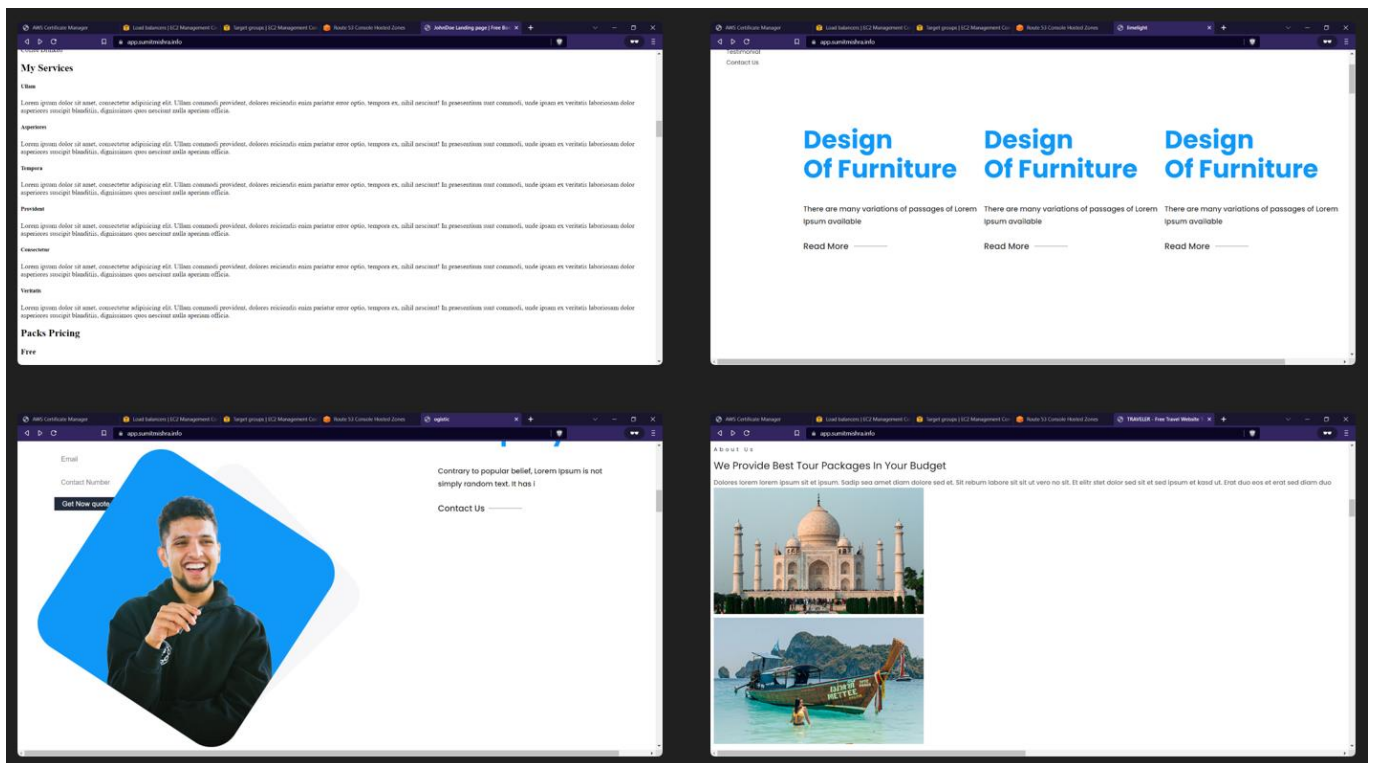
Default SSL/TLS certificate
The certificate used if a client connects without SNI protocol, or if there are no matching certificates. This certificate will automatically be added to your listener certificate list.

From ACM ▼ **sumitmishra.info** ▼ ↻
b5b6c55d-9acd-4417-9390-ce56a147794a

[Request new ACM certificate](#)


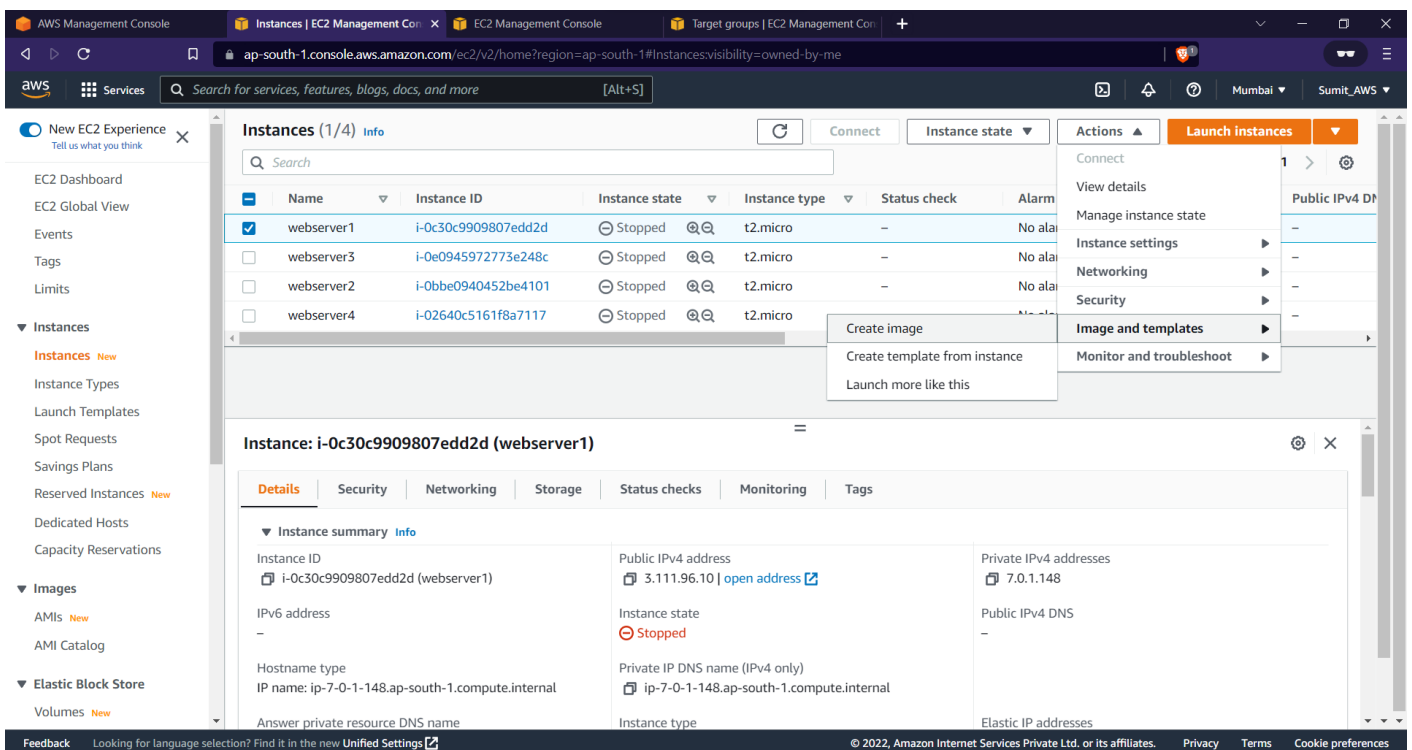
Cancel Save changes

6. SSL has been successfully enabled if we hit the URL with <https://app.sumitmishra.info> and load is balanced between the backend web servers.



5. Configuring auto scaling to an instance.

1. Create an Image of webserver1 EC2 instance.



2. Create a new Load balancer “webserver1-lb” specifically for that server and assign a new configured target group for the “webserver1-tg” not having any instance attached to it.

webserver1-tg

Actions

Details

arn:aws:elasticloadbalancing:ap-south-1:258046353232:targetgroup/webserver1-tg/61b0e35afe4ebf98

Target type Instance	Protocol : Port HTTP: 80	Protocol version HTTP1	VPC vpc-083ef3fe083e0f251
IP address type IPv4	Load balancer None associated		

Total targets	Healthy	Unhealthy	Unused	Initial	Draining
0	0	0	0	0	0

Targets

Monitoring

Health checks

Attributes

Tags

Registered targets (0)

Filter resources by property or value

< 1 >

Instance ID	Name	Port	Zone	Health status	Health status details
-------------	------	------	------	---------------	-----------------------

No registered targets

Register targets

3. Now, created a new launch configuration using image of the “webserver1” instance and also checking on the option of detailed monitoring using AWS cloudwatch.

EC2 > Launch configurations > Create launch configuration

Create launch configuration

Launch configuration name

Name

webserver1-lc

Amazon machine image (AMI)

AMI

webserver1-img

Instance type

Instance type

t2.micro (1 vCPUs, 1 GiB, EBS Only)

Choose instance type

Additional configuration - optional

Purchasing option [Info](#)

☐ Request Spot Instances

IAM instance profile [Info](#)

Select IAM role

Monitoring [Info](#)

☒ Enable EC2 instance detailed monitoring within CloudWatch

EBS-optimized instance

☐ Launch as EBS-optimized instance

► Advanced details

Later, if you want to use a different launch configuration, you can create a new one and apply it to any Auto Scaling group. Existing launch configurations cannot be edited.

Storage (volumes) [Info](#)

EBS volumes

Remove

<input type="checkbox"/>	Volume type	Devices	Snapshot	Size (GiB)	Volume type
<input type="checkbox"/>	Root	/dev/xvda	snap-006011a3351499f7d	8	General purpose SSD (g

+ Add new volume

Free tier eligible customers can get up to 30 GB of EBS storage. [Learn more](#) about free usage tier eligibility and usage restrictions.

Security groups [Info](#)

Assign a security group

☐ Create a new security group

☒ Select an existing security group

Security groups

Copy to new

View rules

Search security groups

< 1 >

<input type="checkbox"/>	Security group ID	Name	VPC ID	Description
<input type="checkbox"/>	sg-05a720fd534a6ba81	personal-vpc-lb-sg	vpc-083ef3fe083e0f251	security group for personal VPC subnet Load balancer
<input type="checkbox"/>	sg-072e8fa63c711a4ff	OpenVPN Access Server-2.8.5-AutogenByAWSMP--1	vpc-01d987e17c3785149	This security group was generated by AWS Marketplace and is based on recommended settings for OpenVPN Access Server version 2.8.5 provided by OpenVPN Inc.
<input type="checkbox"/>	sg-08c979ce29bbae3f5	default	vpc-083ef3fe083e0f251	default VPC security group
<input checked="" type="checkbox"/>	sg-0bd36c936a052ff4e	launch-wizard-3	vpc-083ef3fe083e0f251	launch-wizard-3 created 2022-07-29T09:48:41.971Z
<input type="checkbox"/>	sg-0e9f82d5ed4d73a2d	default	vpc-01d987e17c3785149	default VPC security group

Rules with source of 0.0.0.0/0 allow all IP addresses to access your instance. We recommend setting security group rules to allow access from known IP addresses only.

Key pair (login) [Info](#)

Key pair options

Choose an existing key pair

Existing key pair

VPC-server-sumit

☒ I acknowledge that I have access to the selected private key file (VPC-server-sumit.pem), and that without this file, I won't be able to log into my instance.

Cancel

Create launch configuration

4. Created an auto scaling group using launch dynamic configuration by providing the created launch configuration.

Step 1

Choose launch template or configuration

Step 2

Choose instance launch options

Step 3 (optional)

Configure advanced options

Step 4 (optional)

Configure group size and scaling policies

Step 5 (optional)

Add notifications

Step 6 (optional)

Add tags

Step 7

Review

Choose launch template or configuration [Info](#)

Specify a launch template that contains settings common to all EC2 instances that are launched by this Auto Scaling group. If you currently use launch configurations, you might consider migrating to launch templates.

Name

Auto Scaling group name

Enter a name to identify the group.

webserver1-asg

Must be unique to this account in the current Region and no more than 255 characters.

Launch configuration [Info](#)

[Switch to launch template](#)

Launch configuration

Choose a launch configuration that contains the instance-level settings, such as the Amazon Machine Image (AMI), instance type, key pair, and security groups.

webserver1-lc

[Create a launch configuration](#)

Launch configuration

webserver1-lc

AMI ID

ami-0ced563eea51eed5d

Date created

Mon Aug 01 2022 10:51:52
GMT+0530 (India Standard Time)

Security groups

[sg-0bd36c936a052ff4e](#)

Instance type

t2.micro

Key pair name

VPC-server-sumit

Cancel

Next

Step 1
Choose launch template or configuration

Step 2
Choose instance launch options

Step 3 (optional)
Configure advanced options

Step 4 (optional)
Configure group size and scaling policies

Step 5 (optional)
Add notifications

Step 6 (optional)
Add tags

Step 7
Review

Choose instance launch options [Info](#)

Choose the VPC network environment that your instances are launched into, and customize the instance types and purchase options.

Network [Info](#)

For most applications, you can use multiple Availability Zones and let EC2 Auto Scaling balance your instances across the zones. The default VPC and default subnets are suitable for getting started quickly.

VPC

Choose the VPC that defines the virtual network for your Auto Scaling group.

vpc-083ef3fe083e0f251 (personal-vpc)
7.0.0.0/16



[Create a VPC](#)

Availability Zones and subnets

Define which Availability Zones and subnets your Auto Scaling group can use in the chosen VPC.

Select Availability Zones and subnets



ap-south-1a | subnet-0c89d2c7eab4ba5ac (web-subnet-1)
7.0.1.0/24



ap-south-1b | subnet-0225018fd4692efec (web-subnet-2)
7.0.2.0/24



[Create a subnet](#)

Cancel

Previous

Skip to review

Next

EC2 > Auto Scaling groups > Create Auto Scaling group

Step 1
Choose launch template or configuration

Step 2
Choose instance launch options

Step 3 (optional)
Configure advanced options

Step 4 (optional)
Configure group size and scaling policies

Step 5 (optional)
Add notifications

Step 6 (optional)
Add tags

Step 7
Review

Configure advanced options [Info](#)

Choose a load balancer to distribute incoming traffic for your application across instances to make it more reliable and easily scalable. You can also set options that give you more control over health check replacements and monitoring.

Load balancing - optional [Info](#)

Use the options below to attach your Auto Scaling group to an existing load balancer, or to a new load balancer that you define.

☐ No load balancer
Traffic to your Auto Scaling group will not be fronted by a load balancer.

☒ Attach to an existing load balancer
Choose from your existing load balancers.

☐ Attach to a new load balancer
Quickly create a basic load balancer to attach to your Auto Scaling group.

Attach to an existing load balancer

Select the load balancers that you want to attach to your Auto Scaling group.

☒ Choose from your load balancer target groups
This option allows you to attach Application, Network, or Gateway Load Balancers.

☐ Choose from Classic Load Balancers

Existing load balancer target groups

Only instance target groups that belong to the same VPC as your Auto Scaling group are available for selection.

Select target groups



webserver1-tg | HTTP
Application Load Balancer: webserver1-lb



Health checks - *optional*

Health check type [Info](#)

EC2 Auto Scaling automatically replaces instances that fail health checks. If you enabled load balancing, you can enable ELB health checks in addition to the EC2 health checks that are always enabled.

☒ EC2 ☐ ELB

Health check grace period

The amount of time until EC2 Auto Scaling performs the first health check on new instances after they are put into service.

20 seconds

Additional settings - *optional*

Monitoring [Info](#)

☒ Enable group metrics collection within CloudWatch

Default instance warmup [Info](#)

The amount of time that CloudWatch metrics for new instances do not contribute to the group's aggregated instance metrics, as their usage data is not reliable yet.

☐ Enable default instance warmup

Cancel

Previous

Skip to review

Next

EC2 > Auto Scaling groups > Create Auto Scaling group

Step 1

Choose launch template or configuration

Step 2

Choose instance launch options

Step 3 (optional)

Configure advanced options

Step 4 (optional)

Configure group size and scaling policies

Step 5 (optional)

Add notifications

Step 6 (optional)

Add tags

Step 7

Review

Configure group size and scaling policies [Info](#)

Set the desired, minimum, and maximum capacity of your Auto Scaling group. You can optionally add a scaling policy to dynamically scale the number of instances in the group.

Group size - *optional* [Info](#)

Specify the size of the Auto Scaling group by changing the desired capacity. You can also specify minimum and maximum capacity limits. Your desired capacity must be within the limit range.

Desired capacity

1

Minimum capacity

1

Maximum capacity

3

Scaling policies - *optional*

Choose whether to use a scaling policy to dynamically resize your Auto Scaling group to meet changes in demand. [Info](#)



Target tracking scaling policy

Choose a desired outcome and leave it to the scaling policy to add and remove capacity as needed to achieve that outcome.



None

Instance scale-in protection - *optional*

Instance scale-in protection

If protect from scale in is enabled, newly launched instances will be protected from scale in by default.

☐ Enable instance scale-in protection

Cancel

Previous

Skip to review

Next

EC2 > Auto Scaling groups > Create Auto Scaling group

Step 1
Choose launch template or configuration

Step 2
Choose instance launch options

Step 3 (optional)
Configure advanced options

Step 4 (optional)
Configure group size and scaling policies

Step 5 (optional)
Add notifications

Step 6 (optional)
Add tags

Step 7
Review

Add tags Info

Add tags to help you search, filter, and track your Auto Scaling group across AWS. You can also choose to automatically add these tags to instances when they are launched.


Tags (1)

Key	Value - optional	Tag new instances	
<input type="text" value="Name"/>	<input type="text" value="webserver1-scaled"/>	<input checked="" type="checkbox"/>	<input type="button" value="Remove"/>

49 remaining

When auto scaling group will be created then by default an instance using the image created will be launched in the backend.

Instances (2) Info

<input type="checkbox"/>	Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone
<input type="checkbox"/>	webserver1	i-0c30c9909807edd2d	Running	t2.micro	2/2 checks passed	No alarms +	ap-south-1a
	webserver1-sc...	i-0a038cf7d341fe7d7	Running	t2.micro	2/2 checks passed	No alarms +	ap-south-1a

5. Added a record in route table of my domain for the web server load balancer.

Route 53 > Hosted zones > sumitmishra.info > Create record

Quick create record Info Switch to wizard

Record 1

Record name Info: .sumitmishra.info

Record type Info:

Keep blank to create a record for the root domain.

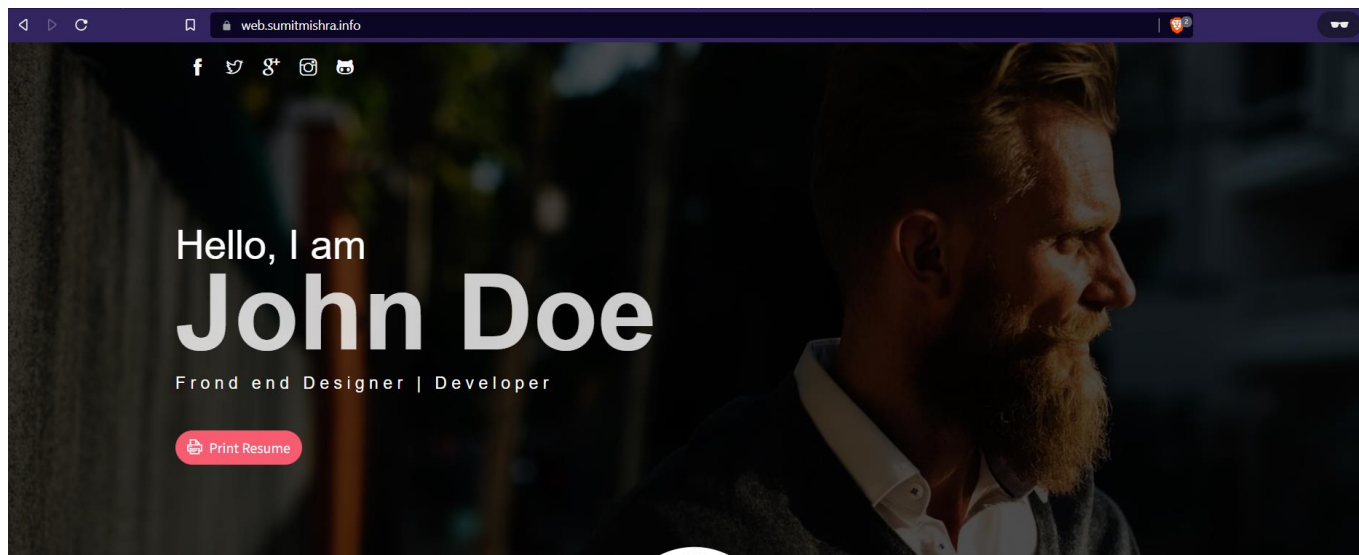
☒ Alias

Route traffic to Info:

Routing policy Info:

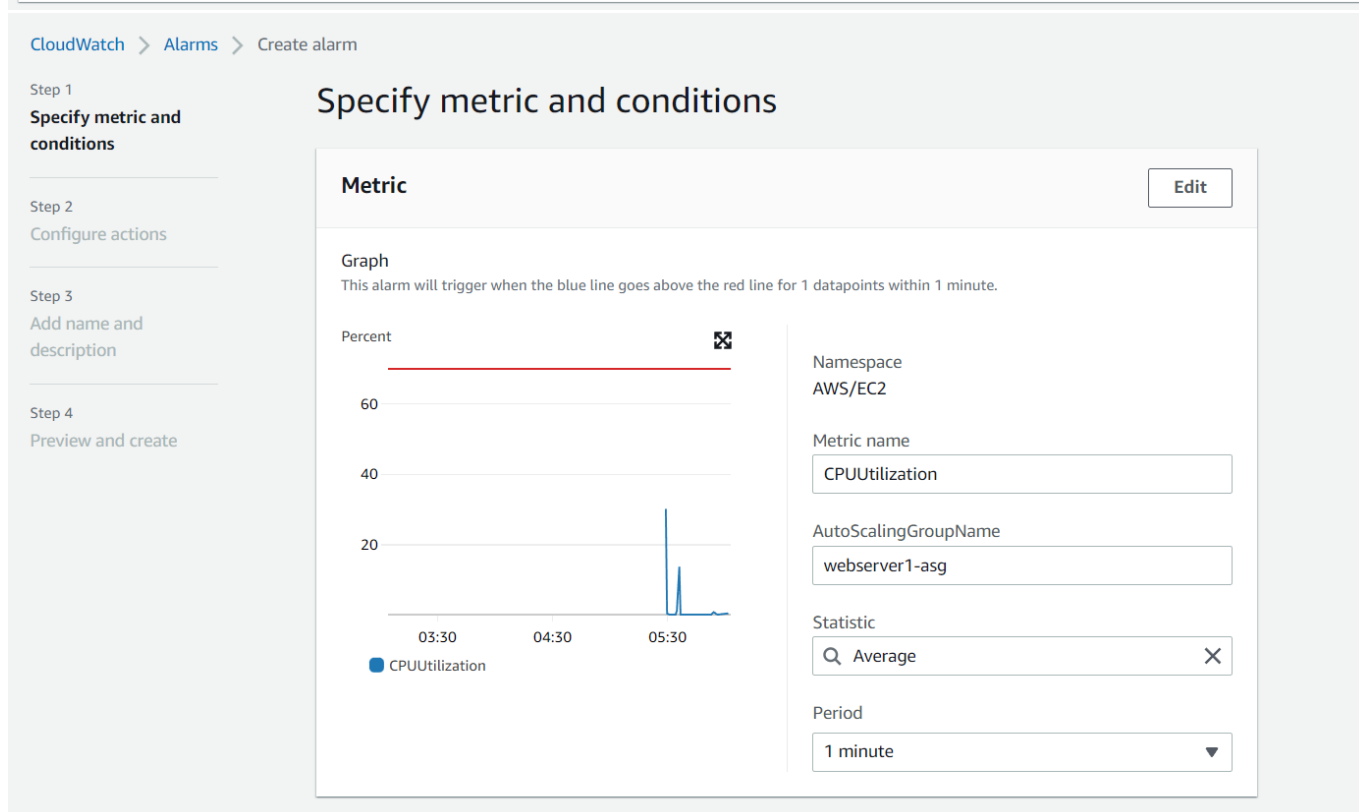
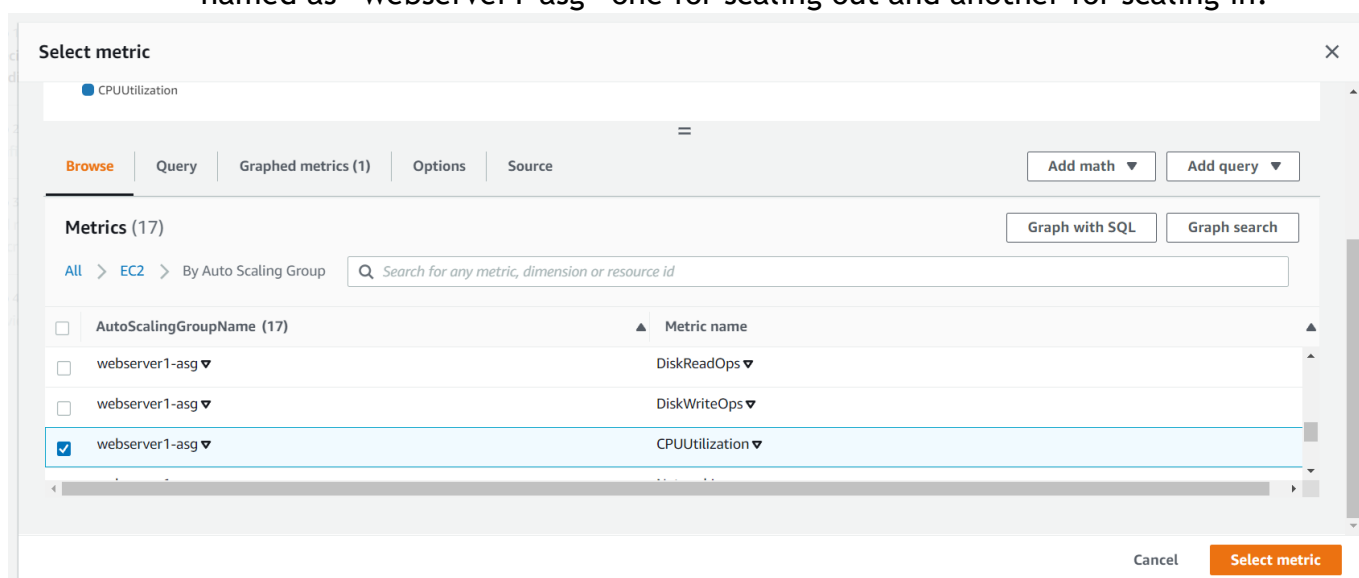
Evaluate target health: ☒ Yes

6. Then I created an SSL certificate for the webserver using AWS certificate manager in the same way as I did for “personal-vpc-lb” (mentioned earlier).



As we can see here the SSL certificate is enabled for <https://web.sumitmishra.info> .

- Created two Cloud watch alarms using the auto scaling group I just created named as “webserver1-asg” one for scaling out and another for scaling in.



For scale-out

Conditions

Threshold type

☒ Static
Use a value as a threshold

☐ Anomaly detection
Use a band as a threshold

Whenever CPUUtilization is...
Define the alarm condition.

☒ Greater
> threshold

☐ Greater/Equal
≥ threshold

☐ Lower/Equal
≤ threshold

☐ Lower
< threshold

than...
Define the threshold value.

70

Must be a number

For scale in

Conditions

Threshold type

☒ Static
Use a value as a threshold

☐ Anomaly detection
Use a band as a threshold

Whenever CPUUtilization is...
Define the alarm condition.

☐ Greater
> threshold

☐ Greater/Equal
≥ threshold

☐ Lower/Equal
≤ threshold

☒ Lower
< threshold

than...
Define the threshold value.

30

Must be a number

► Additional configuration

CancelNext

Alarms created :

CloudWatch > Alarms

Alarms (2) ☐ Hide Auto Scaling alarms Clear selection Refresh Create composite alarm Actions Create alarm

Any state Any type Any actions ... < 1 > Settings

<input type="checkbox"/>	Name	State	Last state update	Conditions	Actions
<input type="checkbox"/>	Scale-In	⚠ In alarm	2022-08-01 11:48:28	CPUUtilization < 30 for 1 datapoints within 1 minute	No actions
<input type="checkbox"/>	Scale-Out	✅ OK	2022-08-01 11:41:47	CPUUtilization > 70 for 1 datapoints within 1 minute	No actions

8. Then I created two dynamic scaling policies for scaling in and scaling out operation, and attached the alarms to respective scaling policies.

For scale-in:

EC2 > Auto Scaling groups > webserver1-asg

Create dynamic scaling policy

Policy type

Simple scaling

Scaling policy name

scale-in

CloudWatch alarm

Choose an alarm that can scale capacity whenever:

Scale-In

[Create a CloudWatch alarm](#)

breaches the alarm threshold: CPUUtilization < 30 for 1 consecutive periods of 60 seconds for the metric dimensions:

AutoScalingGroupName = webserver1-asg

Take the action

Remove

1

capacity units

And then wait

20 seconds before allowing another scaling activity

Cancel Create

For scale-out :

Edit dynamic scaling policy

Policy type

Simple scaling ▼

Scaling policy name

scale-out

CloudWatch alarm

Choose an alarm that can scale capacity whenever:

Scale-Out ▼


[Create a CloudWatch alarm](#)

breaches the alarm threshold: CPUUtilization > 70 for 1 consecutive periods of 60 seconds for the metric dimensions:

AutoScalingGroupName = webserver1-asg

Take the action

Add ▼

1

capacity units ▼

And then wait

20

seconds before allowing another scaling activity

Cancel

Update

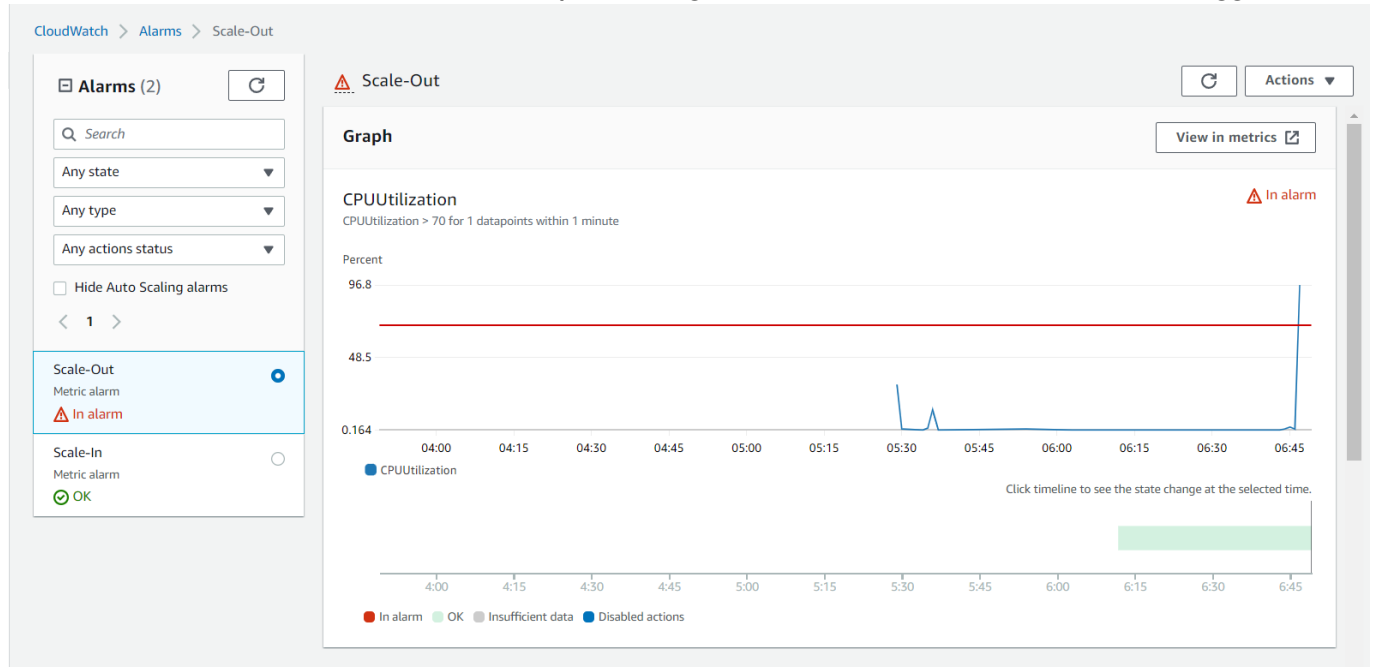
9. Then I increased the CPU utilization to 99% by typing the command “yes /dev/null &” then I typed “top” to see the utilization table.

```
[root@ip-7-0-1-148 ~]# yes > /dev/null &
[1] 3878
[root@ip-7-0-1-148 ~]# top
top - 06:23:59 up 1:57, 1 user, load average: 0.66, 0.24, 0.09
Tasks: 108 total, 2 running, 64 sleeping, 0 stopped, 0 zombie
%Cpu(s): 98.7 us, 1.3 sy, 0.0 ni, 0.0 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
KiB Mem : 988688 total, 616948 free, 112272 used, 259468 buff/cache
KiB Swap: 0 total, 0 free, 0 used. 729984 avail Mem
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
3878	root	20	0	114644	760	696	R	99.7	0.1	0:33.75	yes
2945	apache	20	0	555948	8348	5260	S	0.3	0.8	0:00.94	httpd
2948	apache	20	0	326376	7100	4136	S	0.3	0.7	0:00.81	httpd
1	root	20	0	41588	5264	3816	S	0.0	0.5	0:02.43	systemd
2	root	20	0	0	0	0	S	0.0	0.0	0:00.00	kthreadd
3	root	0	-20	0	0	0	I	0.0	0.0	0:00.00	rcu_gp
4	root	0	-20	0	0	0	I	0.0	0.0	0:00.00	rcu_par_gp
6	root	0	-20	0	0	0	I	0.0	0.0	0:00.00	kworker/0:0H-ev
8	root	0	-20	0	0	0	I	0.0	0.0	0:00.11	kworker/0:1H-ev
9	root	0	-20	0	0	0	I	0.0	0.0	0:00.00	mm_percpu_wq
10	root	20	0	0	0	0	S	0.0	0.0	0:00.00	rcu_tasks_rude_
11	root	20	0	0	0	0	S	0.0	0.0	0:00.00	rcu_tasks_trace
12	root	20	0	0	0	0	S	0.0	0.0	0:00.03	ksoftirqd/0
13	root	20	0	0	0	0	I	0.0	0.0	0:00.09	rcu_sched
14	root	rt	0	0	0	0	S	0.0	0.0	0:00.03	migration/0
16	root	20	0	0	0	0	S	0.0	0.0	0:00.00	cpuhp/0
18	root	20	0	0	0	0	S	0.0	0.0	0:00.00	kdevtmpfs
19	root	0	-20	0	0	0	I	0.0	0.0	0:00.00	netns
22	root	20	0	0	0	0	S	0.0	0.0	0:00.01	kauditd
264	root	20	0	0	0	0	S	0.0	0.0	0:00.00	khungtaskd
265	root	20	0	0	0	0	S	0.0	0.0	0:00.00	oom_reaper
266	root	0	-20	0	0	0	I	0.0	0.0	0:00.00	writeback
268	root	20	0	0	0	0	S	0.0	0.0	0:00.16	kcompactd0
269	root	25	5	0	0	0	S	0.0	0.0	0:00.00	ksmd
270	root	39	19	0	0	0	S	0.0	0.0	0:00.00	khugepaged
325	root	0	-20	0	0	0	I	0.0	0.0	0:00.00	kintegrityd
327	root	0	-20	0	0	0	I	0.0	0.0	0:00.00	kblockd
328	root	0	-20	0	0	0	I	0.0	0.0	0:00.00	blkcg_punt_bio
680	root	20	0	0	0	0	S	0.0	0.0	0:00.00	xen-balloon
686	root	0	-20	0	0	0	I	0.0	0.0	0:00.00	tpm_dev_wq
692	root	0	-20	0	0	0	I	0.0	0.0	0:00.00	md
695	root	0	-20	0	0	0	I	0.0	0.0	0:00.00	edac-poller
700	root	-51	0	0	0	0	S	0.0	0.0	0:00.00	watchdogd
849	root	20	0	0	0	0	S	0.0	0.0	0:00.00	kswapd0

```
ssh://ec2-user@3.111.96.10:22
```

Here, we can see that CPU utilization percentage is 99.7% and scale-out alarm is triggered.



10. In the backend, servers started to deploy automatically to manage the traffic of CPU.

Instances (1/6) Info										Launch Instances
Search										
	Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IPv4 DNS	Public IP	
<input type="checkbox"/>	webserver1	i-0c30c9909807edd2d	Running	t2.micro	2/2 checks passed	No alarms	ap-south-1a	-	3.111.96	
<input type="checkbox"/>	webserver3	i-0e0945972773e248c	Stopped	t2.micro	-	No alarms	ap-south-1a	-	43.205.4	
<input type="checkbox"/>	webserver1-sc...	i-0a038cf7d341fe7d7	Running	t2.micro	2/2 checks passed	No alarms	ap-south-1a	-	13.127.1	
<input type="checkbox"/>	webserver2	i-0bbe0940452be4101	Stopped	t2.micro	-	No alarms	ap-south-1b	-	3.109.30	
<input type="checkbox"/>	webserver4	i-02640c5161f8a7117	Stopped	t2.micro	-	No alarms	ap-south-1b	-	43.205.6	
<input checked="" type="checkbox"/>	webserver1-sc...	i-0a461f60bcd96504	Running	t2.micro	Initializing	No alarms	ap-south-1b	-	13.234.6	

11. After that I decreased the CPU utilization of the server using the command “killall -p yes”. Then used “top” command to get the metrics.

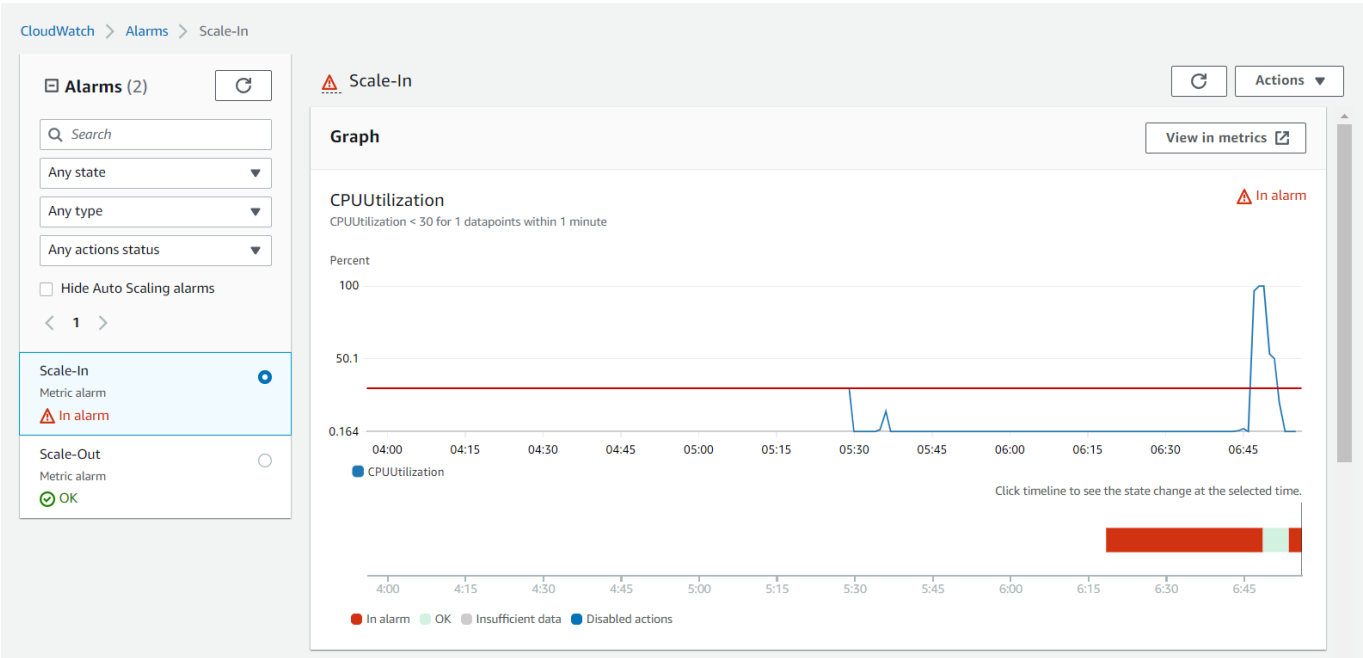
We can see that CPU utilization has dropped close to 0%.

```
[root@ip-7-0-1-206 ~]# killall -p yes
[root@ip-7-0-1-206 ~]# top
```

```
top - 06:53:09 up 1:24, 1 user, load average: 0.47, 0.57, 0.29
Tasks: 106 total, 1 running, 64 sleeping, 0 stopped, 0 zombie
%Cpu(s): 0.0 us, 0.0 sy, 0.0 ni, 100.0 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
KiB Mem : 988688 total, 613868 free, 112408 used, 262412 buff/cache
KiB Swap: 0 total, 0 free, 0 used. 729572 avail Mem

  PID USER      PR  NI    VIRT    RES    SHR  S  %CPU  %MEM     TIME+ COMMAND
    1 root        20   0 123504    5396   3892  S   0.0   0.5   0:02.11 systemd
    2 root        20   0        0        0        0  S   0.0   0.0   0:00.00 kthreadd
    3 root         0 -20        0        0        0  I   0.0   0.0   0:00.00 rcu_gp
    4 root         0 -20        0        0        0  I   0.0   0.0   0:00.00 rcu_par_gp
    6 root         0 -20        0        0        0  I   0.0   0.0   0:00.00 kworker/0:0H-ev
    8 root         0 -20        0        0        0  I   0.0   0.0   0:00.07 kworker/0:1H-ev
    9 root         0 -20        0        0        0  I   0.0   0.0   0:00.00 mm_percpu_wq
   10 root        20   0        0        0        0  S   0.0   0.0   0:00.00 rcu_tasks_rude_
   11 root        20   0        0        0        0  S   0.0   0.0   0:00.00 rcu_tasks_trace
   12 root        20   0        0        0        0  S   0.0   0.0   0:00.03 ksoftirqd/0
   13 root        20   0        0        0        0  I   0.0   0.0   0:00.11 rcu_sched
   14 root        rt    0        0        0        0  S   0.0   0.0   0:00.02 migration/0
   16 root        20   0        0        0        0  S   0.0   0.0   0:00.00 cpuhp/0
   18 root        20   0        0        0        0  S   0.0   0.0   0:00.00 kdevtmpfs
   19 root         0 -20        0        0        0  I   0.0   0.0   0:00.00 netns
   20 root        20   0        0        0        0  I   0.0   0.0   0:00.02 kworker/u30:1-e
   22 root        20   0        0        0        0  S   0.0   0.0   0:00.01 kauditd
  264 root        20   0        0        0        0  S   0.0   0.0   0:00.00 khungtaskd
  265 root        20   0        0        0        0  S   0.0   0.0   0:00.00 oom_reaper
  266 root         0 -20        0        0        0  I   0.0   0.0   0:00.00 writeback
  268 root        20   0        0        0        0  S   0.0   0.0   0:00.10 kcompactd0
  269 root        25   5        0        0        0  S   0.0   0.0   0:00.00 ksmd
  270 root        39  19        0        0        0  S   0.0   0.0   0:00.00 khugepaged
  325 root         0 -20        0        0        0  I   0.0   0.0   0:00.00 kintegrityd
  327 root         0 -20        0        0        0  I   0.0   0.0   0:00.00 kblockd
```

We can see that the scale in alarm has been triggered in the alarm section.



12. In the backend, the allocated servers started to terminate.

Instances (1/6) Info

Search

< 1 >

	Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IPv4 DNS	Public IP
<input type="checkbox"/>	webserver1	i-0c30c9909807edd2d	● Running	t2.micro	● 2/2 checks passed	No alarms	ap-south-1a	-	3.111.96.10
<input type="checkbox"/>	webserver3	i-0e0945972773e248c	● Stopped	t2.micro	-	No alarms	ap-south-1a	-	43.205.14.10
<input checked="" type="checkbox"/>	webserver1-sc...	i-0a038cf7d341fe7d7	● Shutting-down	t2.micro	-	No alarms	ap-south-1a	-	13.127.14.10
<input type="checkbox"/>	webserver2	i-0bbe0940452be4101	● Stopped	t2.micro	-	No alarms	ap-south-1b	-	3.109.30.10
<input type="checkbox"/>	webserver4	i-02640c5161f8a7117	● Stopped	t2.micro	-	No alarms	ap-south-1b	-	43.205.14.10
<input type="checkbox"/>	webserver1-sc...	i-0a461f60bcd96504	● Running	t2.micro	● 2/2 checks passed	No alarms	ap-south-1b	-	13.234.14.10